

# Title of Project

## ✓ BANK CUSTOMER CHURN MODEL

### ✓ Objective

1. Data Encoding

2. Feature Scaling

3. Handling Imbalance Data

1. Random Under Sampling

2. Random Over Sampling

4. Support Vector Machine Classifier

5. Grid Search for Hyperparameter Tunning

### ✓ Data Source

<https://github.com/YBI-Foundation/Dataset>

Double-click (or enter) to edit

### ✓ Import Library

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

### ✓ Import Data

```
df = pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/Bank%20Churn%20Modellir
```

## ✓ Analyze Data

```
df.head()
```



	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products
0	15634602	Hargrave	619	France	Female	42	2	0.00	1
1	15647311	Hill	608	Spain	Female	41	1	83807.86	1
2	15619304	Onio	502	France	Female	42	8	159660.80	3
3	15701354	Boni	699	France	Female	39	1	0.00	2

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerId            10000 non-null  int64
1   Surname                10000 non-null  object
2   CreditScore            10000 non-null  int64
3   Geography              10000 non-null  object
4   Gender                 10000 non-null  object
5   Age                    10000 non-null  int64
6   Tenure                 10000 non-null  int64
7   Balance                 10000 non-null  float64
8   Num Of Products        10000 non-null  int64
9   Has Credit Card        10000 non-null  int64
10  Is Active Member       10000 non-null  int64
11  Estimated Salary       10000 non-null  float64
12  Churn                   10000 non-null  int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB
```

```
df.duplicated('CustomerId').sum()
```



```
0
```

```
df = df.set_index('CustomerId')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 10000 entries, 15634602 to 15628319
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Surname                10000 non-null  object
 1   CreditScore            10000 non-null  int64
 2   Geography              10000 non-null  object
 3   Gender                 10000 non-null  object
 4   Age                   10000 non-null  int64
 5   Tenure                 10000 non-null  int64
 6   Balance                10000 non-null  float64
 7   Num Of Products        10000 non-null  int64
 8   Has Credit Card        10000 non-null  int64
 9   Is Active Member       10000 non-null  int64
10   Estimated Salary       10000 non-null  float64
11   Churn                   10000 non-null  int64
dtypes: float64(2), int64(7), object(3)
memory usage: 1015.6+ KB
```

## ✓ Data Encoding

```
df['Geography'].value_counts()
```

```
Geography
France      5014
Germany     2509
Spain       2477
Name: count, dtype: int64
```

```
df.replace({'Geography' : {'France' :2, 'Germany': 1, 'Spain': 0}}, inplace=True)
```

```
df['Gender'].value_counts()
```

```
Gender
Male      5457
Female    4543
Name: count, dtype: int64
```

```
df.replace({'Gender': {'Male': 0, 'Female':1}}, inplace=True)
```

```
df['Num Of Products'].value_counts()
```

```
Num Of Products
1      5084
2      4590
3       266
```

```
4      60
Name: count, dtype: int64
```

```
df.replace({'Num Of Products': {1:0, 2:1, 3:1, 4:1}}, inplace=True)
```

```
df['Has Credit Card'].value_counts()
```

```
⇒ Has Credit Card
1      7055
0      2945
Name: count, dtype: int64
```

```
df['Is Active Member'].value_counts()
```

```
⇒ Is Active Member
1      5151
0      4849
Name: count, dtype: int64
```


```
df.loc[(df['Balance']==0), 'Churn'].value_counts()
```

```
⇒ Churn
0      3117
1        500
Name: count, dtype: int64
```

```
df['Zero Balance'] = np.where(df['Balance']>0,1,0)
```

```
df['Zero Balance'].hist()
```

↗ <Axes: >



```
df.groupby(['Churn', 'Geography']).count()
```

↗

		Surname	CreditScore	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card
Churn	Geography								
0	0	2064	2064	2064	2064	2064	2064	2064	2064
	1	1695	1695	1695	1695	1695	1695	1695	1695
	2	4204	4204	4204	4204	4204	4204	4204	4204
1	0	413	413	413	413	413	413	413	413
	1	814	814	814	814	814	814	814	814

## ✓ Define Label and Features

```
df.columns
```

```
↗ Index(['Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
        'Balance', 'Num Of Products', 'Has Credit Card', 'Is Active Member',
        'Estimated Salary', 'Churn', 'Zero Balance'],
        dtype='object')
```

```
x = df.drop(['Surname', 'Churn'], axis=1)
```

```
y = df['Churn']
```

```
x.shape,y.shape
```


```
↵↵ ((10000, 11), (10000,))
```

## ✓ Undersampling and Oversampling

```
df['Churn'].value_counts()
```

```
↵↵ Churn  
0    7963  
1    2037  
Name: count, dtype: int64
```

```
sns.countplot(x='Churn', data = df);
```



```
x.shape,y.shape
```

```
↵↵ ((10000, 11), (10000,))
```

## ✓ Random Under Sampling

```
from imblearn.under_sampling import RandomUnderSampler
```

```
rus = RandomUnderSampler(random_state=2529)
```

```
x_rus, y_rus = rus.fit_resample(x,y)
```

```
x_rus.shape, y_rus.shape, x.shape, y.shape
```

```
→ ((4074, 11), (4074,), (10000, 11), (10000,))
```

```
y.value_counts()
```


```
→ Churn  
0    7963  
1    2037  
Name: count, dtype: int64
```

```
y_rus.value_counts()
```

```
→ Churn  
0    2037  
1    2037  
Name: count, dtype: int64
```

```
y_rus.plot(kind = 'hist')
```

↩ <Axes: ylabel='Frequency'>



## ✓ Random Over Sampling

```
from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state=2529)
```

```
x_ros, y_ros = ros.fit_resample(x,y)
```

```
x_ros.shape, y_ros.shape, x.shape, y.shape
```

↩ ((15926, 11), (15926,), (10000, 11), (10000,))

```
y.value_counts()
```

↩ Churn

0	7963
1	2037

Name: count, dtype: int64


```
y_ros.value_counts()
```



```
Churn
1    7963
0    7963
Name: count, dtype: int64
```

```
y_ros.plot(kind = 'hist')
```

```
<Axes: ylabel='Frequency'>
```



## ✓ Train Test Split

```
from sklearn.model_selection import train_test_split
```

## ✓ Split Original Data

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3, random_state=2529)
```

## ✓ Split Random Under Sample Data

```
x_train_rus,x_test_rus,y_train_rus,y_test_rus = train_test_split(x_rus,y_rus, test_size=0.3,
```

## ✓ Split Random Over Sample Data

```
x_train_ros,x_test_ros,y_train_ros,y_test_ros = train_test_split(x_ros,y_ros, test_size=0.3,
```

## ✓ Standardize Features

Double-click (or enter) to edit

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

## ✓ Standardize Original Data

```
x_train[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_tr
```

```
x_test[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_tes
```

## ✓ Standardize Random Under Sample Data

```
x_train_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(
```

```
x_test_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(>
```

## ✓ Standardize Random Over Sample Data

```
x_train_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(
```

```
x_test_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(>
```

## ✓ Support Vector Machine Classifier

```
from sklearn.svm import SVC
```

```
svc = SVC()
```

```
svc.fit(x_train, y_train)
```

```
→ ▾ SVC
   SVC()
```

```
y_pred = svc.predict(x_test)
```

## ✓ Model Accuracy

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
confusion_matrix(y_test, y_pred)
```

```
→ array([[2381,  33],
        [ 436, 150]])
```

```
print(classification_report(y_test, y_pred))
```

```
→
```

	precision	recall	f1-score	support
0	0.85	0.99	0.91	2414
1	0.82	0.26	0.39	586
accuracy			0.84	3000
macro avg	0.83	0.62	0.65	3000
weighted avg	0.84	0.84	0.81	3000

## ✓ Hyperparameter Tuning

```
from sklearn.model_selection import GridSearchCV
```


```
param_grid = {'C': [0.1,1,10],
              'gamma': [1,0.1,0.01],
              'kernel': ['rbf'],
              'class_weight': ['balanced']}
```

```
grid = GridSearchCV(SVC(),param_grid,refit=True,verbose=2, cv=2)
grid.fit(x_train,y_train)
```

```

➡ Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 2.9s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.0s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.6s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.0s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.5s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.9s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.3s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.1s

```



```
print(grid.best_estimator_)
```

```
➡ SVC(C=10, class_weight='balanced', gamma=1)
```

```
grid_predictions = grid.predict(x_test)
```

```
confusion_matrix(y_test,grid_predictions)
```

```
➡ array([[2159, 255],
        [ 343, 243]])
```

```
print(classification_report(y_test,grid_predictions))
```

```

➡
              precision    recall  f1-score   support

     0       0.86         0.89         0.88         2414
     1       0.49         0.41         0.45          586

 accuracy                   0.80         3000
 macro avg                  0.68         0.65         0.66         3000
 weighted avg               0.79         0.80         0.79         3000

```

## ✓ Model with Random Under Sampling

```
svc_rus = SVC()
```

```
svc_rus.fit(x_train_rus, y_train_rus)
```

```
→ SVC
   SVC()
```

```
y_pred_rus = svc_rus.predict(x_test_rus)
```

## ✓ Model Accuracy

```
confusion_matrix(y_test_rus, y_pred_rus)
```

```
→ array([[470, 157],
        [174, 422]])
```

```
print(classification_report(y_test_rus, y_pred_rus))
```

```
→
```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

## ✓ Hyperparameter Tuning

```
param_grid = {'C': [0.1, 1, 10],
              'gamma': [1, 0.1, 0.01],
              'kernel': ['rbf'],
              'class_weight': ['balanced']}
```

```
grid_rus = GridSearchCV(SVC(), param_grid, refit=True, verbose=2, cv=2)
grid_rus.fit(x_train_rus, y_train_rus)
```

```

➞ Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.4s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.3s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s

```

```

└─ GridSearchCV
  └─ estimator: SVC
    └─ SVC

```

```
print(grid_rus.best_estimator_)
```

```
➞ SVC(C=1, class_weight='balanced', gamma=0.1)
```

```
grid_predictions_rus = grid_rus.predict(x_test_rus)
```

```
confusion_matrix(y_test_rus,grid_predictions_rus)
```

```
➞ array([[476, 151],
        [172, 424]])
```

```
print(classification_report(y_test_rus,grid_predictions_rus))
```

```

➞
              precision    recall  f1-score   support

     0       0.73         0.76         0.75         627
     1       0.74         0.71         0.72         596


 accuracy          0.74         1223
 macro avg         0.74         0.74         0.74         1223
 weighted avg      0.74         0.74         0.74         1223

```

## ✓ Model with Random Over Sampling

```
svc_ros = SVC()
```

```
svc_ros.fit(x_train_ros, y_train_ros)
```




▾ SVC  
SVC()

```
y_pred_ros = svc_ros.predict(x_test_ros)
```


## ✓ Model Accuracy

```
confusion_matrix(y_test_ros,y_pred_ros)
```



```
array([[1823,  556],
       [ 626, 1773]])
```

```
print(classification_report(y_test_ros,y_pred_ros))
```



	precision	recall	f1-score	support
0	0.74	0.77	0.76	2379
1	0.76	0.74	0.75	2399
accuracy			0.75	4778
macro avg	0.75	0.75	0.75	4778
weighted avg	0.75	0.75	0.75	4778

## ✓ Hyperparameter Tuning

```
param_grid = {'C': [0.1,1,10],
              'gamma': [1,0.1,0.01],
              'kernel': ['rbf'],
              'class_weight': ['balanced']}
```

```
grid_ros = GridSearchCV(SVC(),param_grid,refit=True,verbose=2, cv=2)
grid_ros.fit(x_train_ros,y_train_ros)
```



Fitting 2 folds for each of 9 candidates, totalling 18 fits

```
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 6.0s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.8s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.9s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.0s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.1s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 4.1s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.1s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.4s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.7s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.7s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 9.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 10.0s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 6.4s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 7.1s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 7.3s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 6.2s
```

▸ **GridSearchCV**

▸ **estimator: SVC**

▸ SVC

```
print(grid_ros.best_estimator_)
```



```
SVC(C=10, class_weight='balanced', gamma=1)
```

```
grid_predictions_ros = grid_ros.predict(x_test_ros)
```

```
confusion_matrix(y_test_ros, grid_predictions_ros)
```

